




Infrastructure as Code with Terraform



```

graph LR
    Business[Business] --> BusinessAnalyst[Business Analyst]
    BusinessAnalyst --> SolutionArchitect[Solution Architect / Technical Lead]
  
```

Business

Business Analyst

Solution Architect
/ Technical Lead



laaC

Imperative

```
ec2.sh

#!/bin/bash

IP_ADDRESS="10.2.2.1"

EC2_INSTANCE=$(ec2-run-instances --instance-type
t2.micro ami-0edab43b6fa892279)

INSTANCE=$(echo ${EC2_INSTANCE} | sed 's/*INSTANCE //'
| sed 's/ .*//')

# Wait for instance to be ready
while ! ec2-describe-instances $INSTANCE | grep -q
"running"
do
    echo Waiting for $INSTANCE is to be ready...
done

# Check if instance is not provisioned and exit
if [ ! $(ec2-describe-instances $INSTANCE | grep -q
"running") ]; then
    echo Instance $INSTANCE is stopped.
    exit
fi

ec2-associate-address $IP_ADDRESS -i $INSTANCE

echo Instance $INSTANCE was created successfully!!!
```

main.tf

```
resource "aws_instance" "webserver" {
    ami          = "ami-0edab43b6fa892279"
    instance_type = "t2.micro"
}
```

Declarative

ec2.yaml

```
- amazon.aws.ec2:
    key_name: mykey
    instance_type: t2.micro
    image: ami-123456
    wait: yes
    group: webserver
    count: 3
    vpc_subnet_id: subnet-29e63245
    assign_public_ip: yes
```





What Is Infrastructure as (from) Code?

- Infrastructure as code (IaC) is an **approach** to infrastructure automation based on practices from software development.
- It emphasizes **consistent, repeatable** routines for **provisioning and changing** systems and their configuration.
- **Changes are made to definitions** and then rolled out to systems through **unattended processes** that include thorough **validation**.

Tooling Categories...



Ad Hoc Scripts

Configuration
Management (CM)
Tools

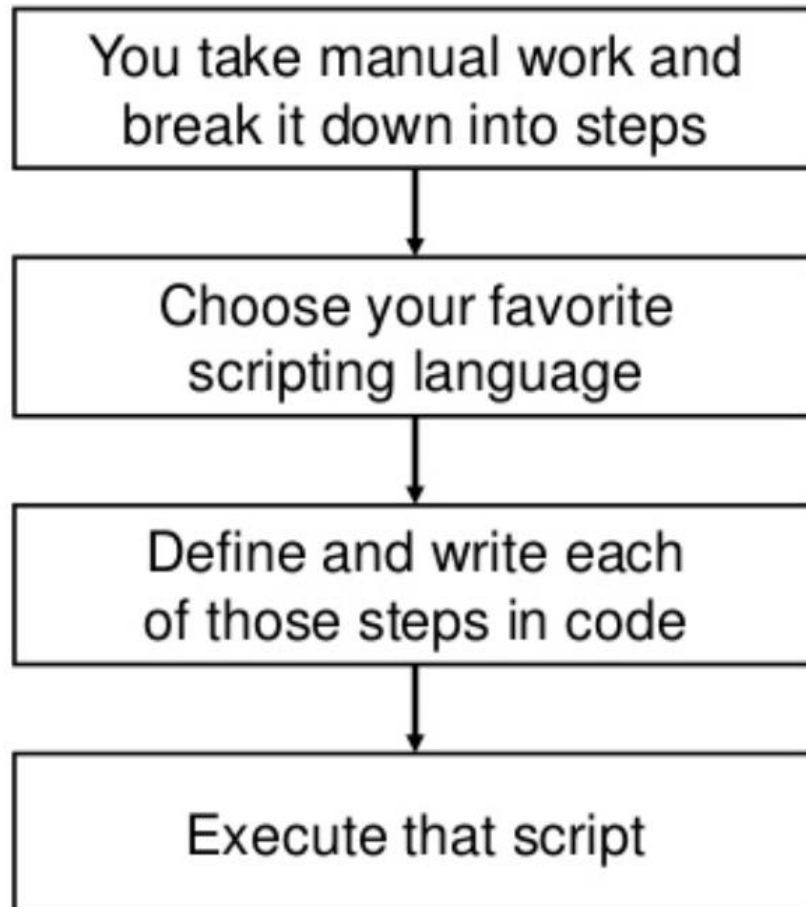
Server Templating
Tools

Server Provisioning
Tools

Ad-Hoc Scripts..



- The most straightforward approach to automating anything is to write an ad hoc script.



```
# Update the apt-get cache
sudo apt-get update

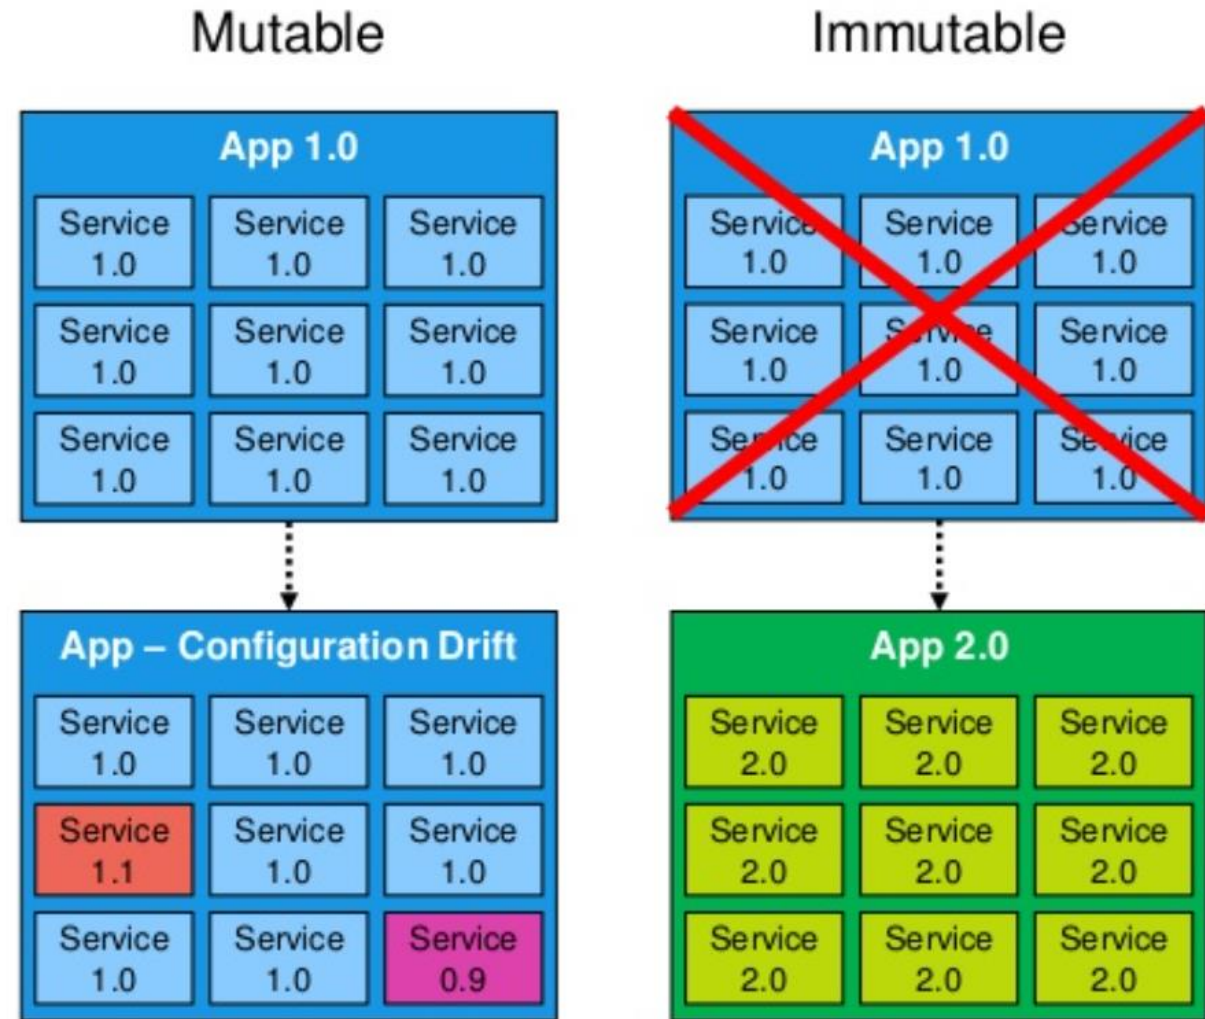
# Install PHP
sudo apt-get install -y php

# Install Apache
sudo apt-get install -y apache2

# Start Apache
sudo service apache2 start
```


Mutable & Immutable Infrastructure

- The Pets and Cattle debate.
- One approach is not necessarily better than the other, it depends on your use-case.
- With the mutable approach, the team needs to be aware of the infrastructure “history”.
- Generally speaking, the immutable approach is better for stateless applications.
- Immutable drives no deviation and no changes. It is what it is.





Infrastructure as Code



SALTSTACK



Types of IAC Tools



Configuration Management



ANSIBLE



SALTSTACK

Install and Manage SW
Version control
Idempotent / Mutable

Server Templating



Pre installed SW with
Dependencies
VM or docker images
Immutable architecture

Provisioning Tools



CloudFormation

Deploy Immutable infr resrce
Servers, DBs, NW components
Multi providers



Terraform is an open-source infrastructure as code software tool that enables you to safely and predictably create, change, and improve infrastructure.











Terraform is an open-source infrastructure as code software tool that provides a consistent CLI workflow to manage hundreds of cloud services. Terraform codifies cloud APIs into declarative configuration files

<https://www.terraform.io/>

HashiCorp was founded by Mitchell Hashimoto and Armon Dadgar in 2012 with the goal of revolutionizing datacenter management: application development, delivery, and maintenance.



Why terraform





Installing Terraform

Download - [Download Terraform - Terraform by HashiCorp](https://learn.hashicorp.com/tutorials/terraform/install-cli)

<https://learn.hashicorp.com/tutorials/terraform/install-cli>

Install yum-config-manager to manage your repositories

```
sudo yum install -y yum-utils
```

#Use yum-config-manager to add the official HashiCorp Linux repository

```
sudo yum-config-manager --add-repo https://rpm.releases.hashicorp.com/AmazonLinux/hashicorp.repo
```

Install terraform

```
sudo yum -y install terraform
```

```
terraform -install-autocomplete
```

HCL



main.tf

```
resource "aws_instance" "webserver" {
  ami          = "ami-0edab43b6fa892279"
  instance_type = "t2.micro"
}

resource "aws_s3_bucket" "finance" {
  bucket = "finanace-21092020"
  tags = {
    Description = "Finance and Payroll"
  }
}

resource "aws_iam_user" "admin-user" {
  name = "lucy"
  tags = {
    Description = "Team Leader"
  }
}
```

Configure the Microsoft Azure Provider

```
provider "azurerm" {
  features {}
}
```

Create a resource group

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West Europe"
}
```

Create a virtual network within the resource group

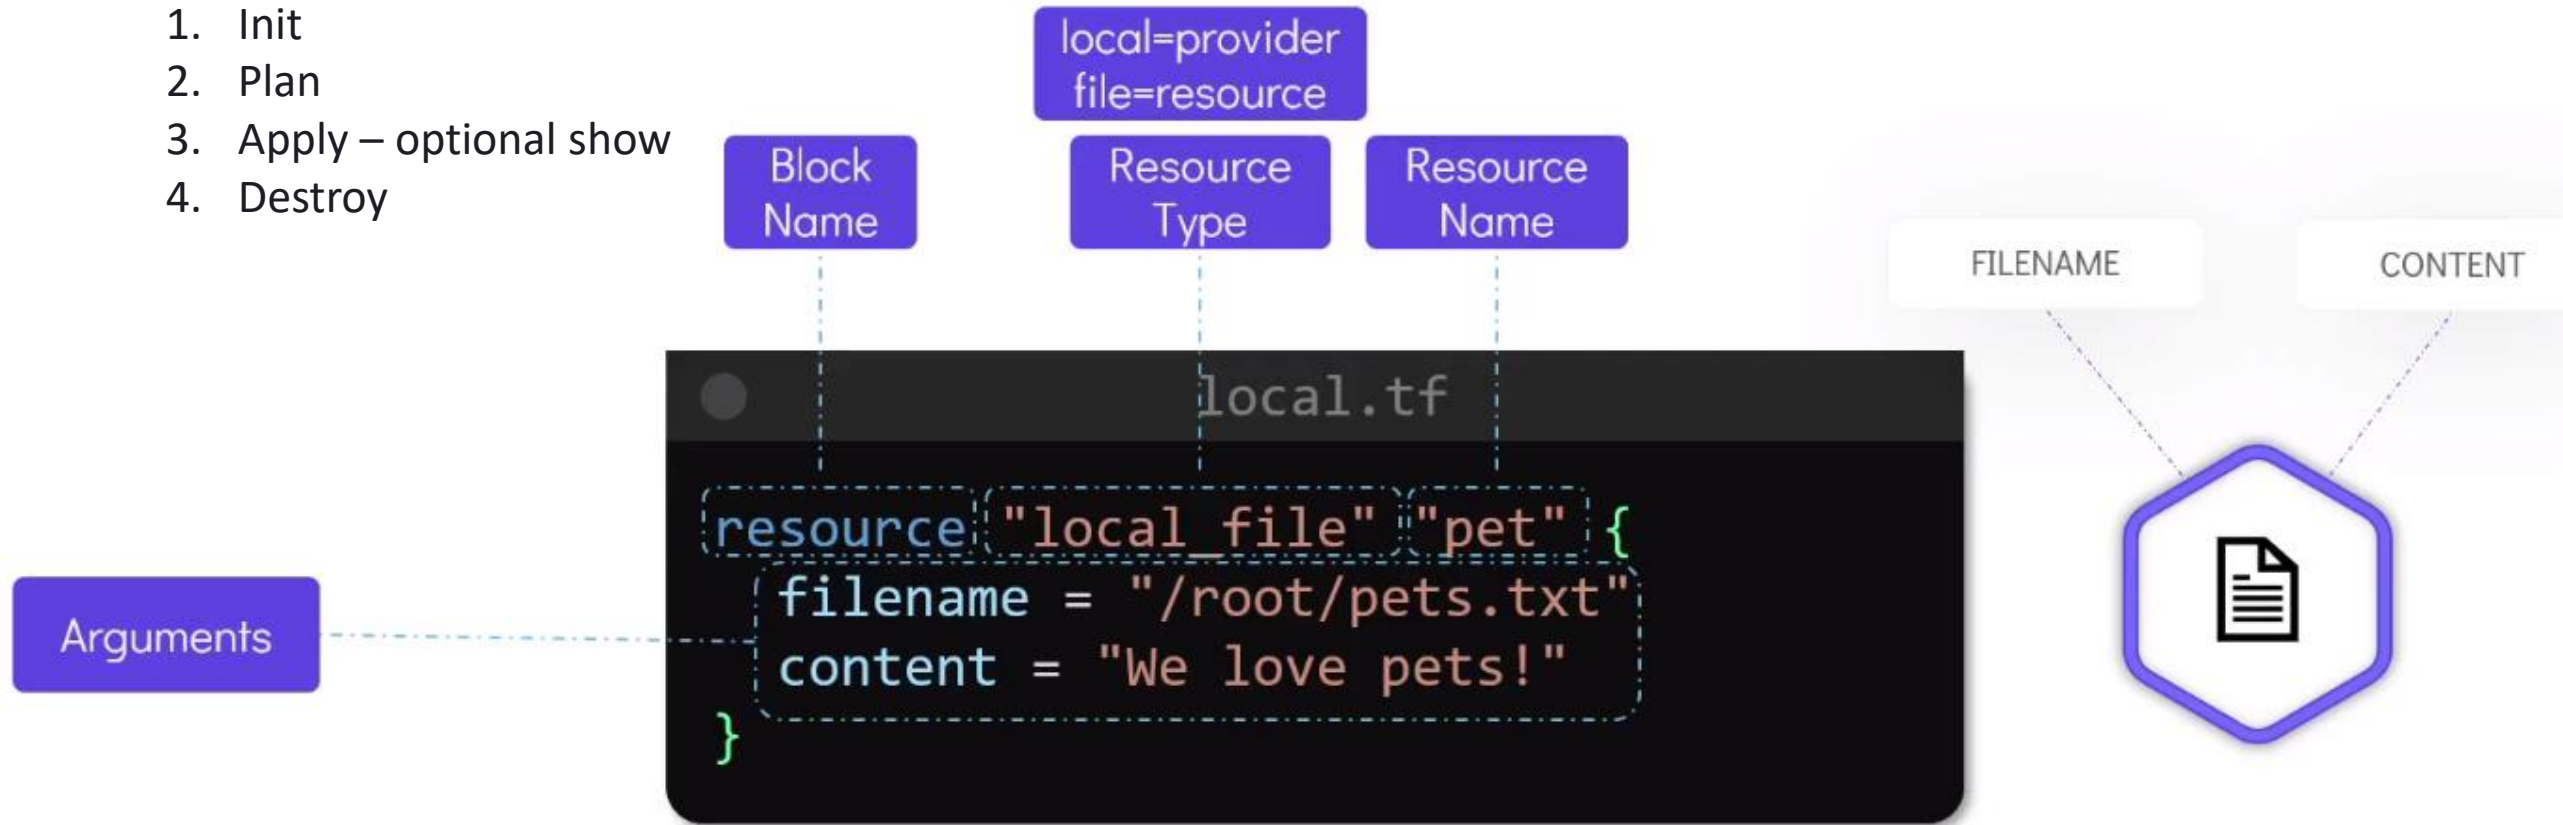
```
resource "azurerm_virtual_network" "example" {
  name                = "example-network"
  resource_group_name = azurerm_resource_group.example.name
  location            = azurerm_resource_group.example.location
  address_space       = ["10.0.0.0/16"]
}
```



.tf



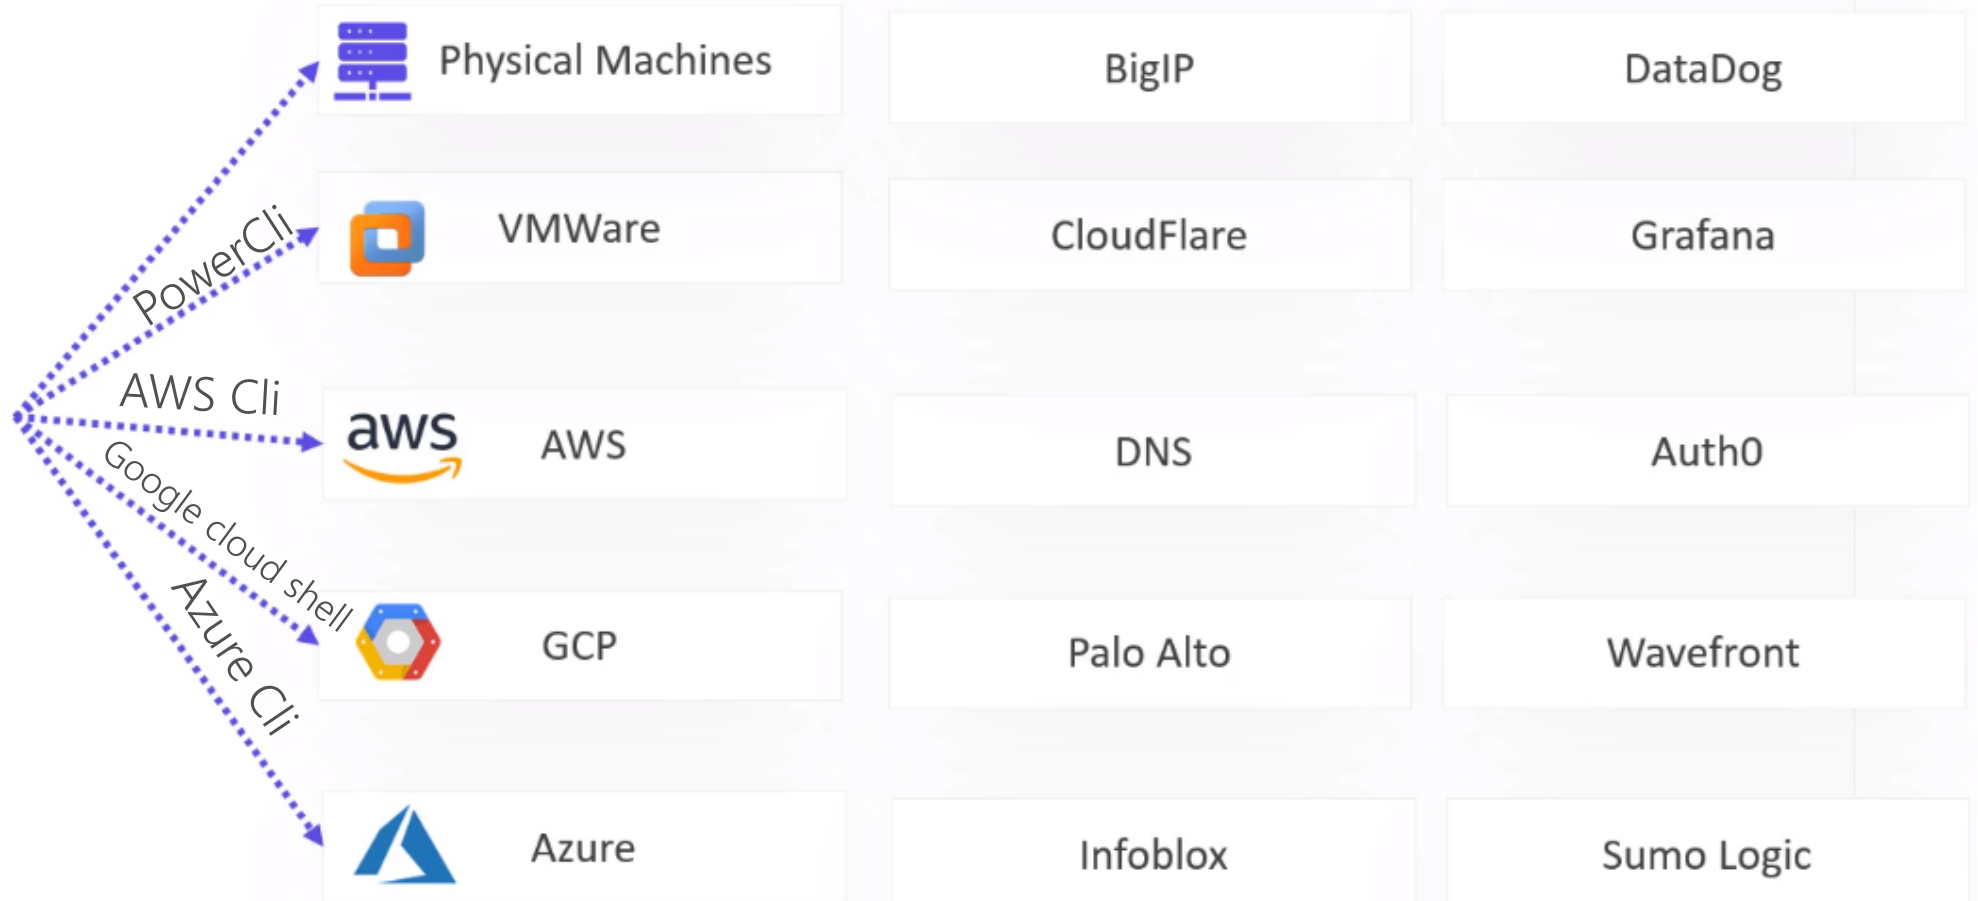
Lets get in to action

1. Create a folder
2. Create a new file ending with .tf extension
3. Terraform commands
 1. Init
 2. Plan
 3. Apply – optional show
 4. Destroy





Why terraform



HCL Resource...




HCL – Declarative Language

```
aws.tf

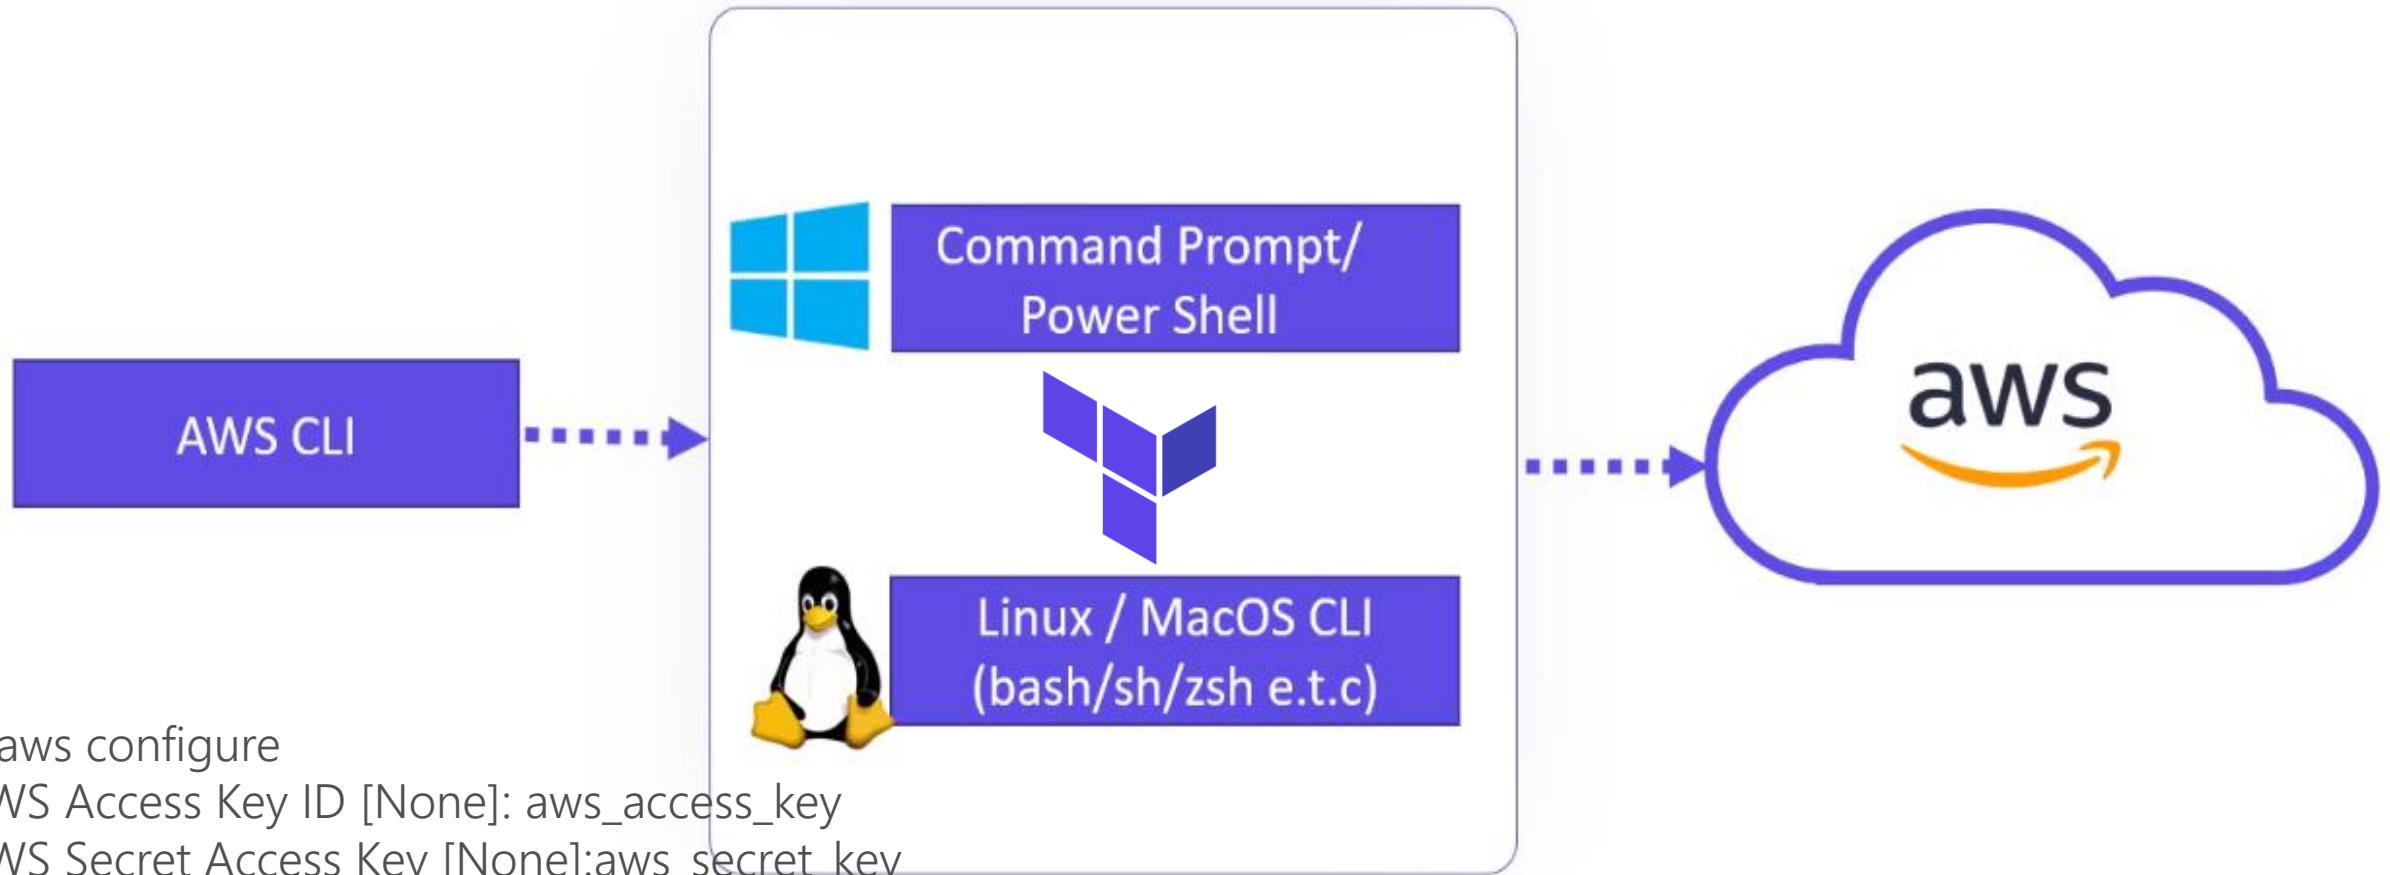
resource "aws_instance" "webserver" {
  ami = "ami-0c2f25c1f66a1ff4d"
  instance_type = "t2.micro"
}
```

The Terraform language or (HCL) is Terraform's primary user interface. In every edition of Terraform, a configuration written in the Terraform language is always at the heart of the workflow.



To manage AWS - Pre-Req

<https://docs.aws.amazon.com/cli/latest/userguide/install-cliv2-linux.html>



```
# aws configure
```

```
AWS Access Key ID [None]: aws_access_key
```

```
AWS Secret Access Key [None]:aws_secret_key
```

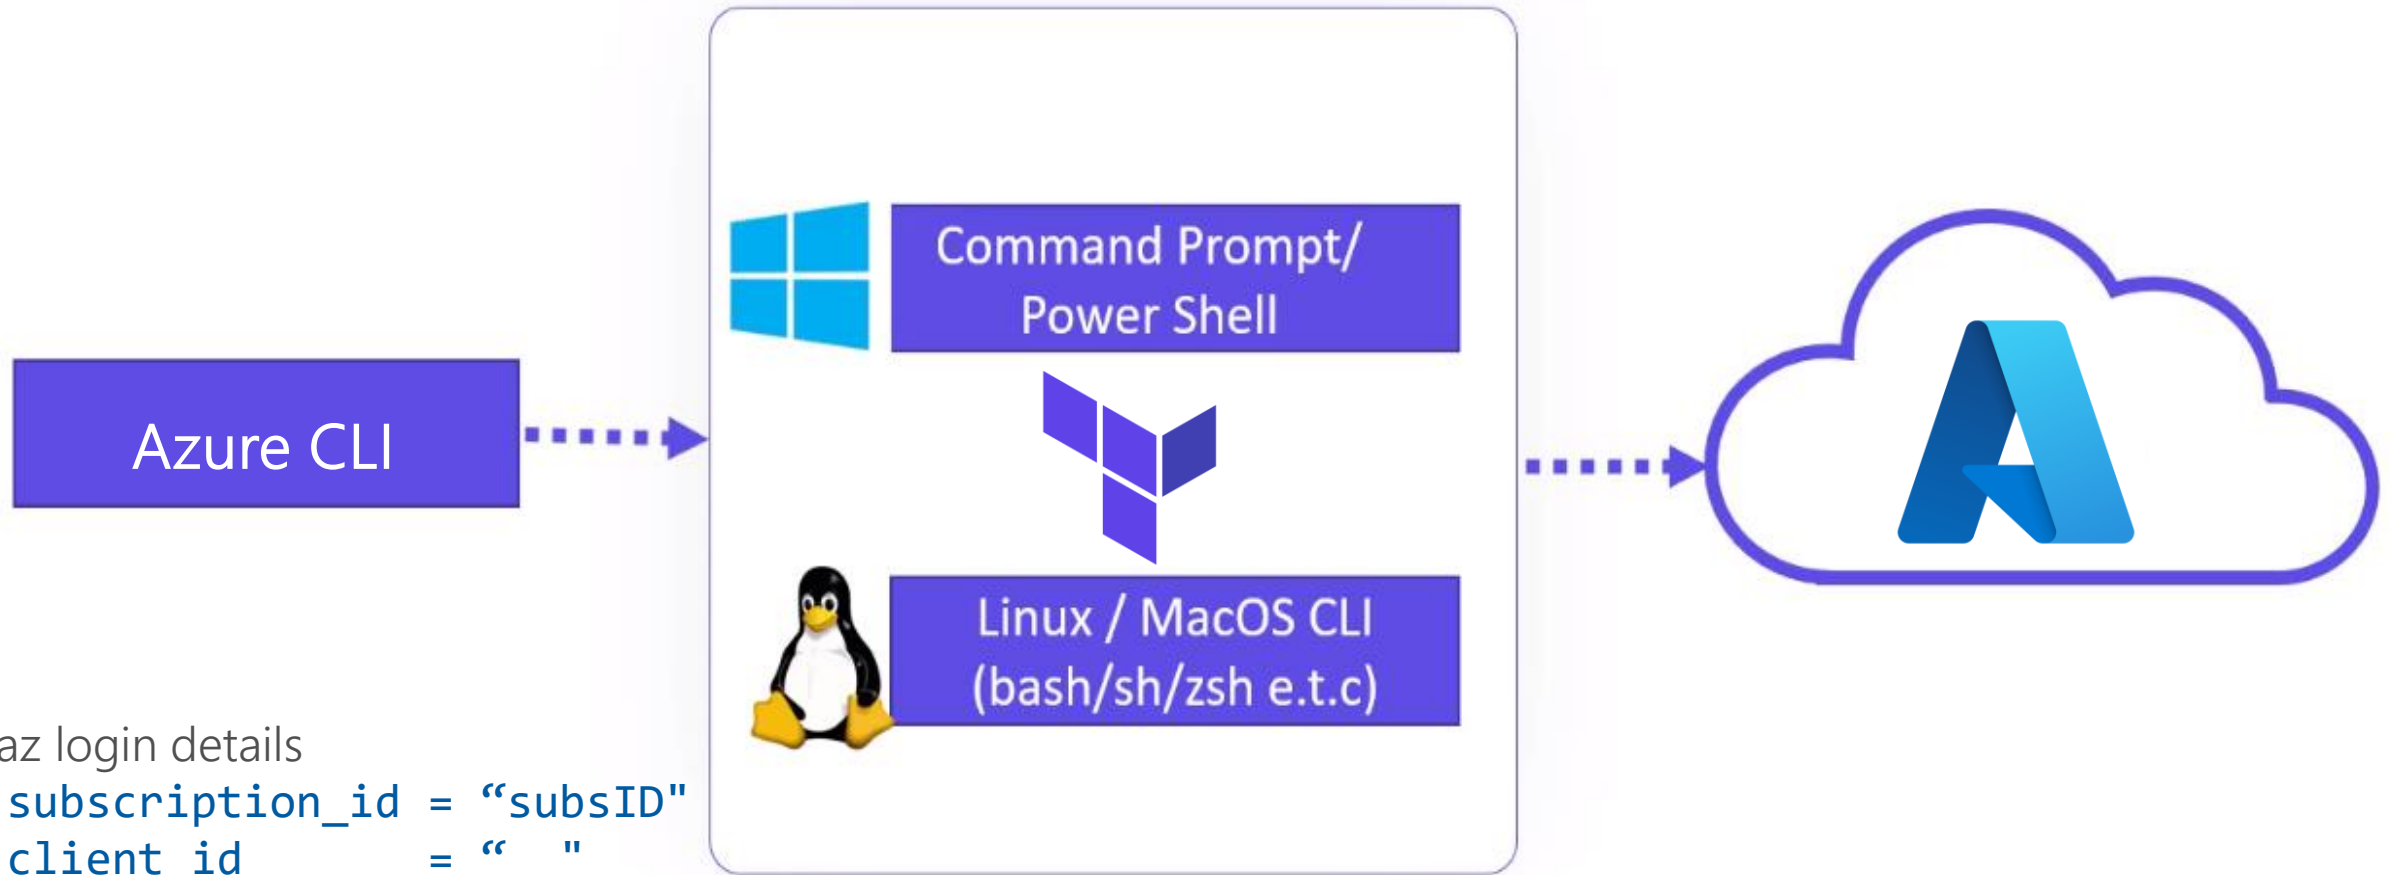
```
Default region name [None]: us-west-2
```

```
Default output format [None]: ENTER
```




To manage Azure - Pre-Req

<https://learn.microsoft.com/en-us/cli/azure/install-azure-cli-linux?pivots=apt>



az login details

```
subscription_id = "subsID"  
client_id       = " "  
client_secret   = " "  
tenant_id      = " "
```

How...



```
# Configure the Microsoft Azure Provider
provider "azurerm" {
  features {}
}

# Create a resource group
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West Europe"
}


# Create a virtual network within the resource group
resource "azurerm_virtual_network" "example" {
  name                 = "example-network"
  resource_group_name =
    azurerm_resource_group.example.name
  location              =
    azurerm_resource_group.example.location
  address_space        = ["10.0.0.0/16"]
}
```

terraform

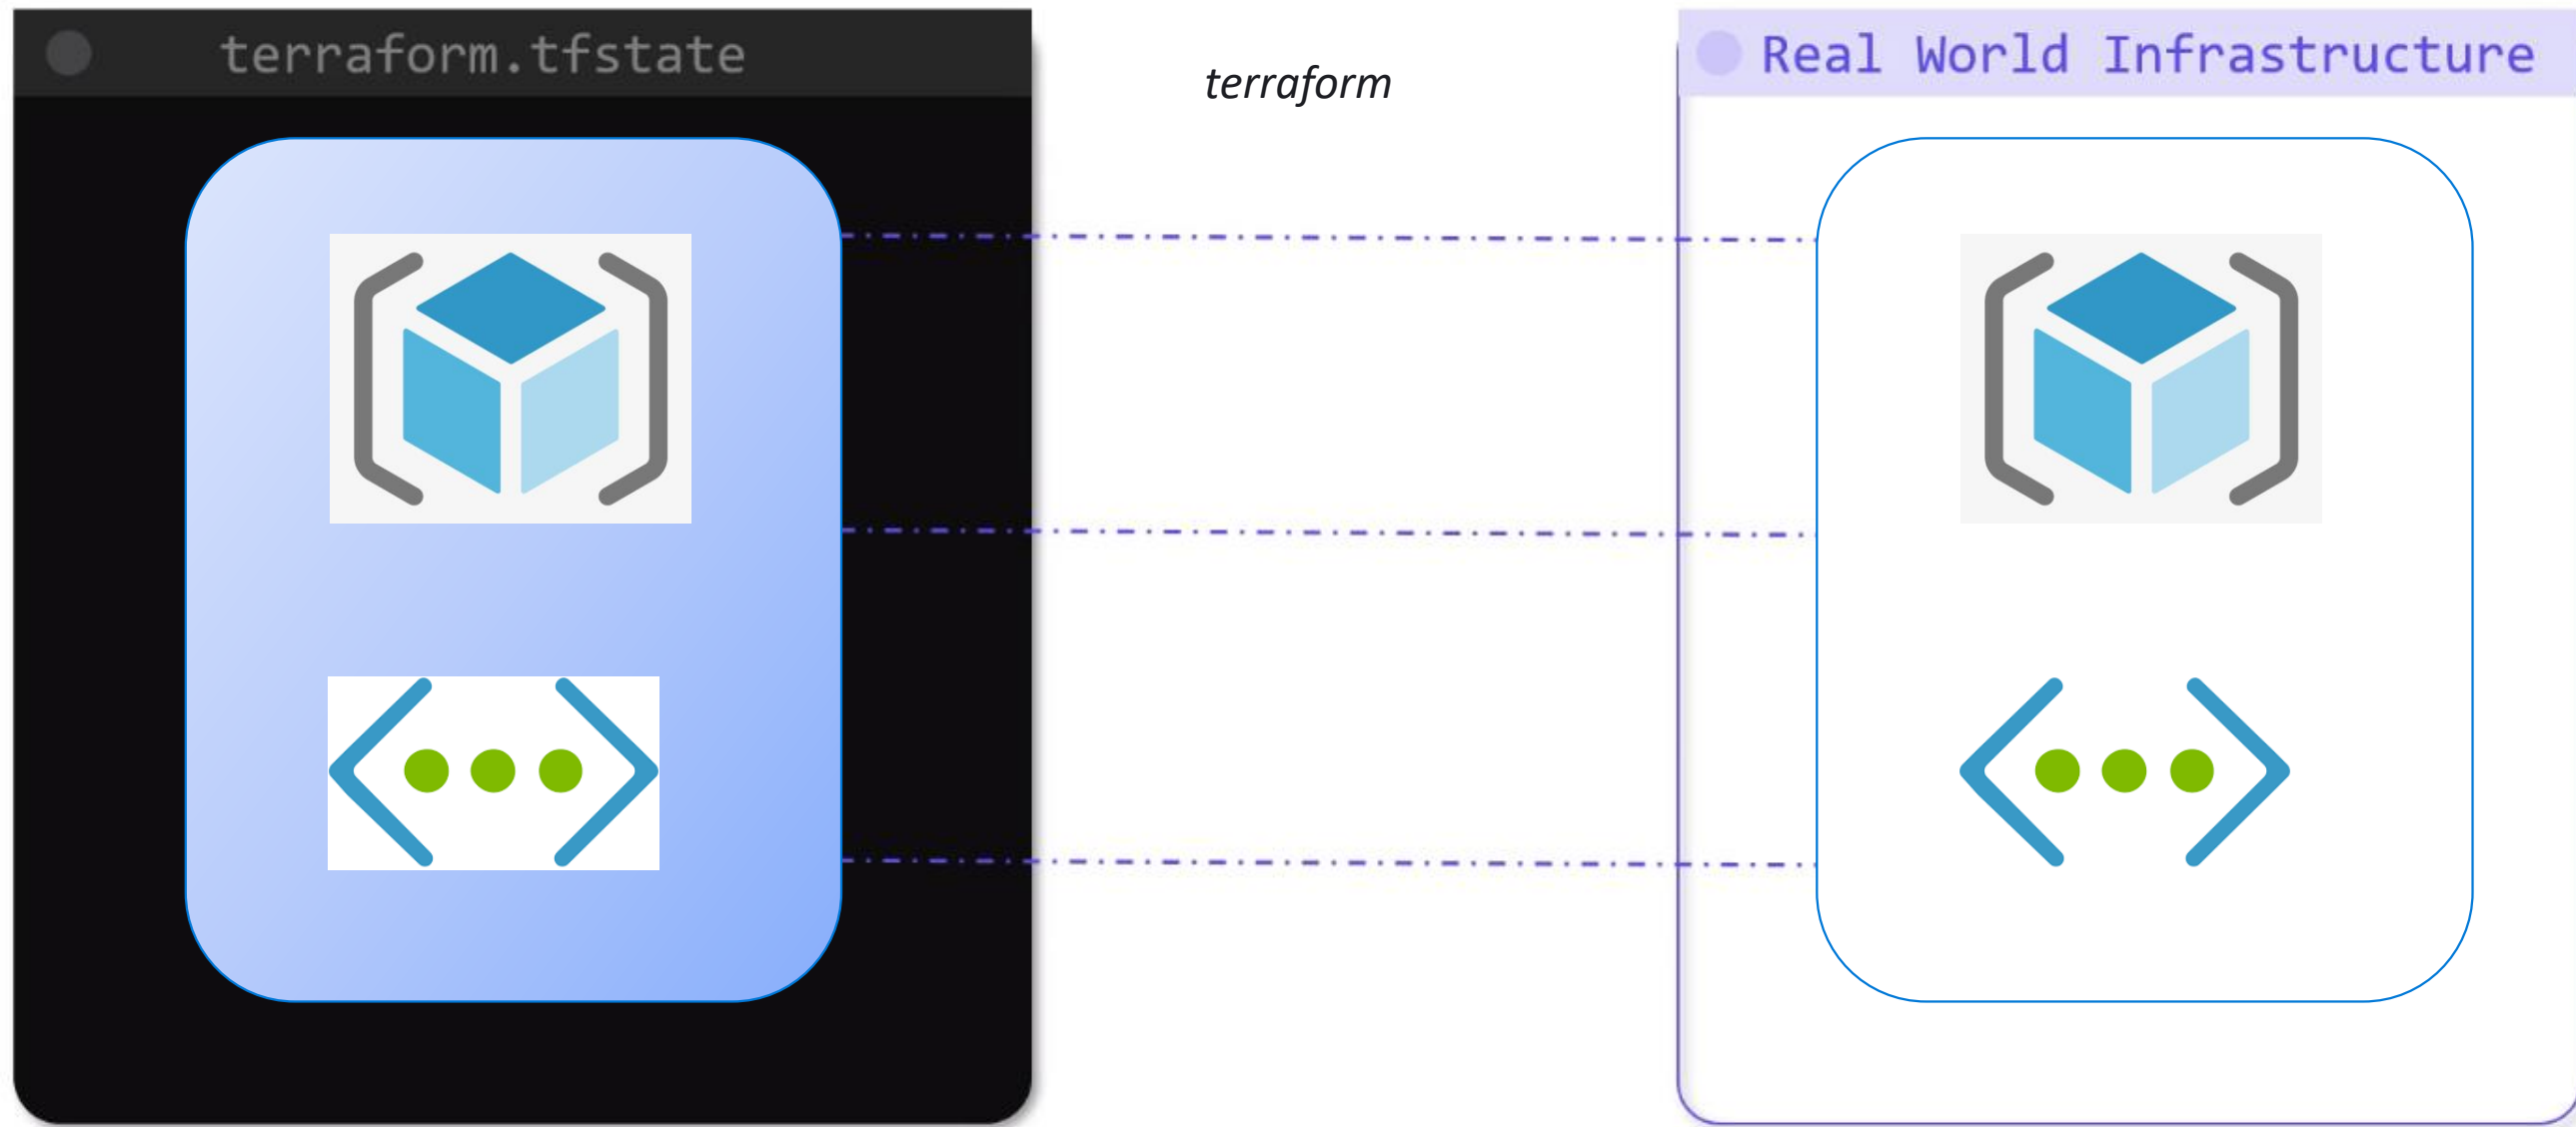
Init

Plan

Apply



Terraform State...



How...



main.tf

```
resource "aws_instance" "webserver" {  
  ami          = "ami-0edab43b6fa892279"  
  instance_type = "t2.micro"  
}  
  
resource "aws_s3_bucket" "finance" {  
  bucket = "finanace-21092020"  
  tags   = {  
    Description = "Finance and Payroll"  
  }  
}  
  
resource "aws_iam_user" "admin-user" {  
  name = "lucy"  
  tags = {  
    Description = "Team Leader"  
  }  
}
```

terraform

Init

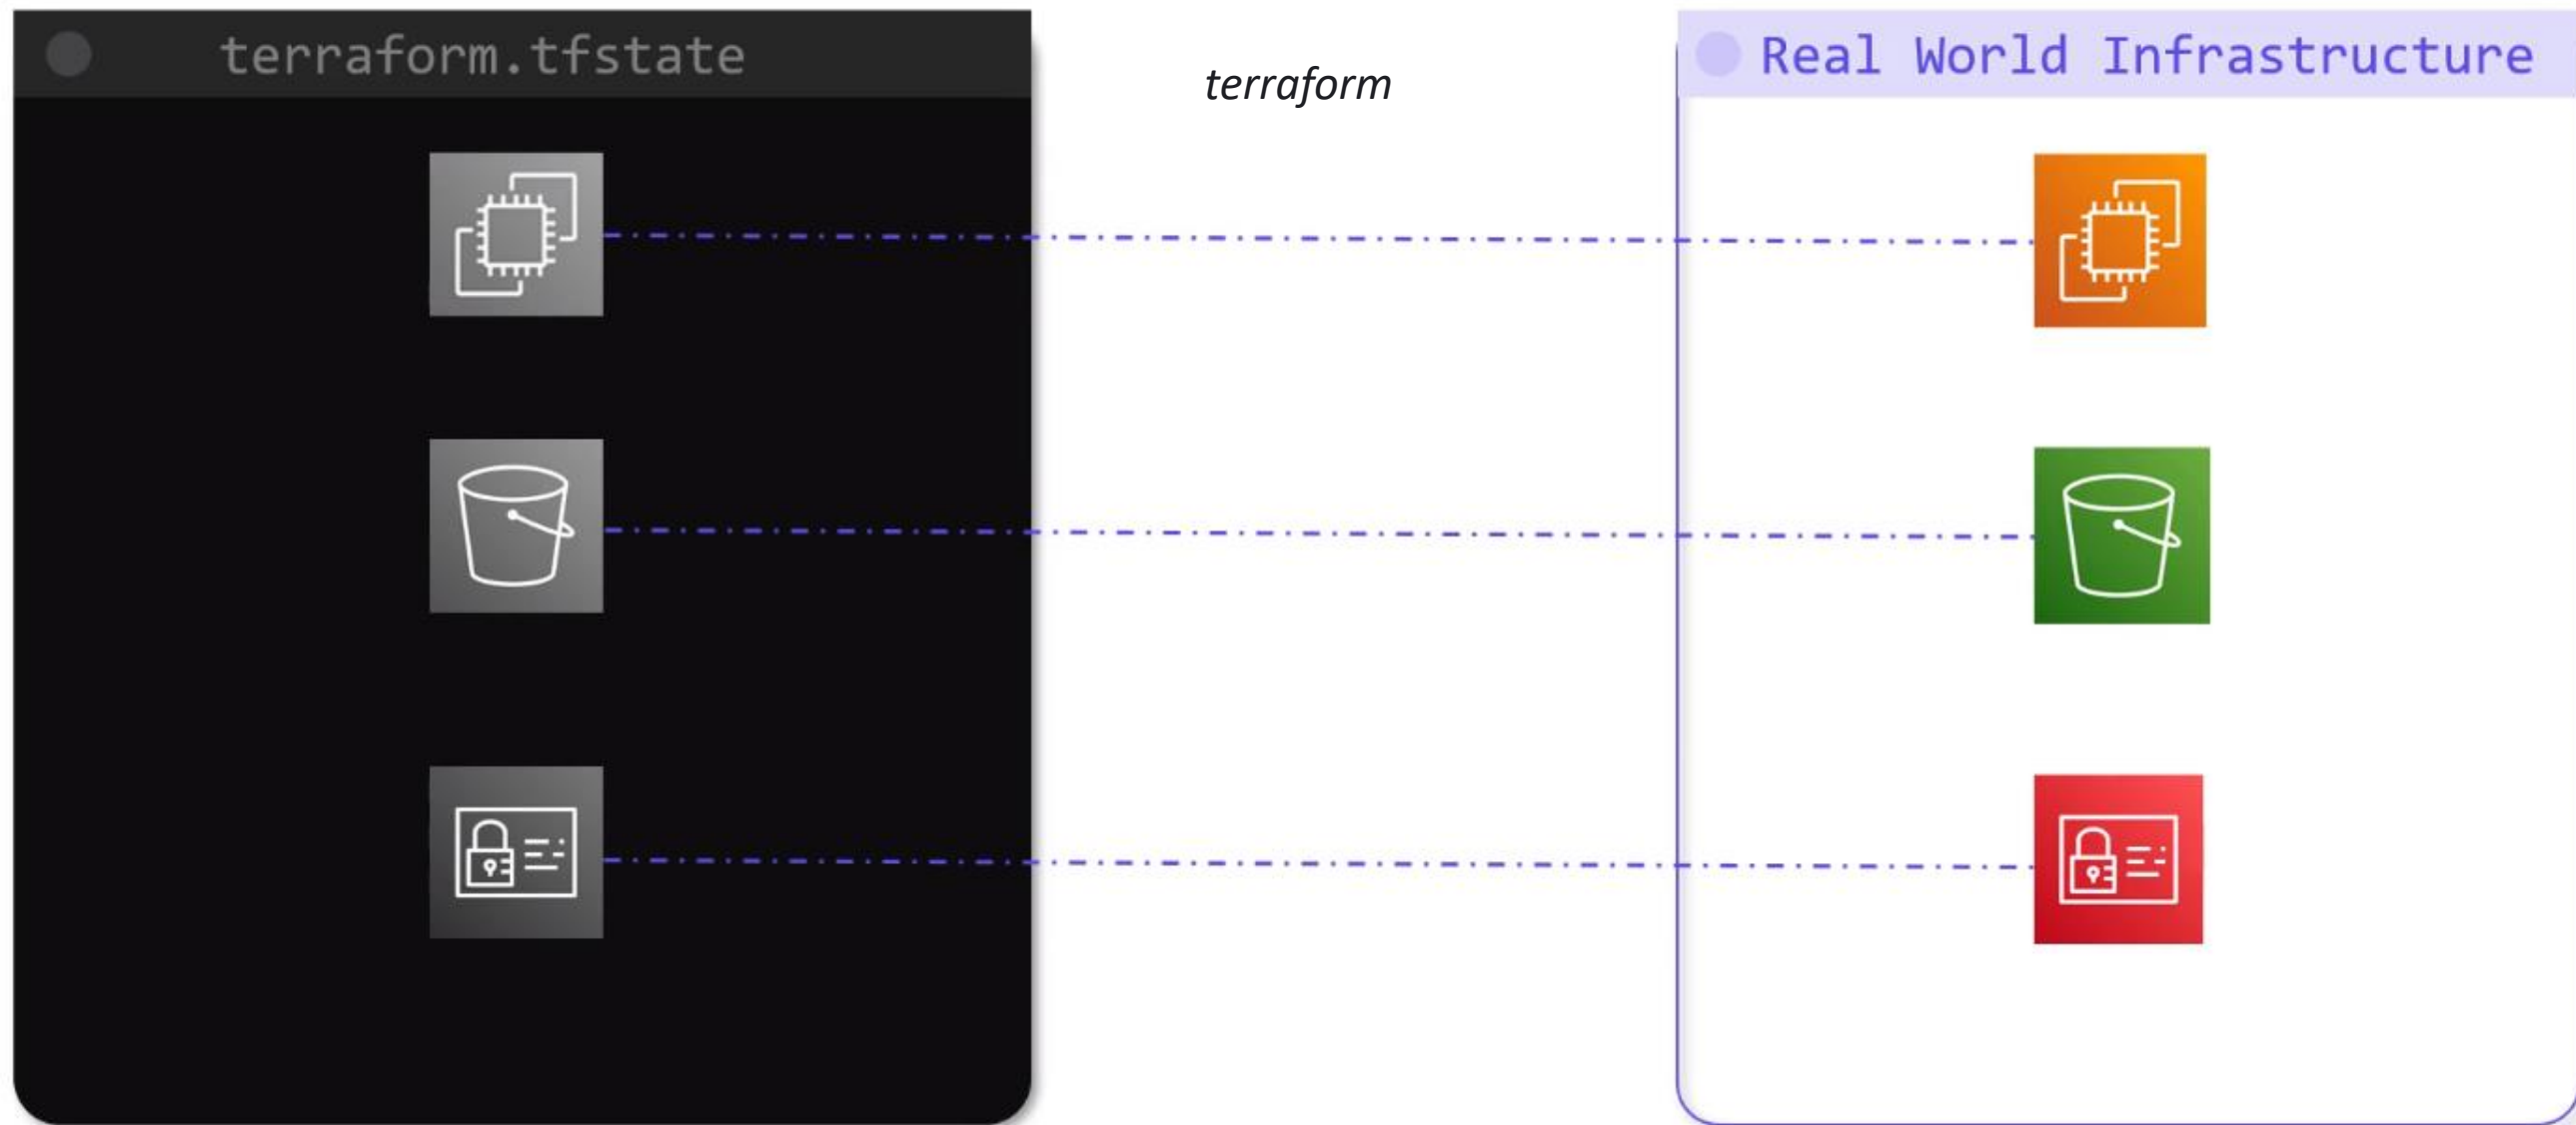
Plan

Apply

Real World Infrastructure



Terraform State...





Terraform Language / HCL Syntax

<https://www.terraform.io/docs/configuration/index.html>

<https://www.terraform.io/docs/configuration/syntax.html>

```
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {  
  # Block body  
  <IDENTIFIER> = <EXPRESSION> # Argument  
}
```

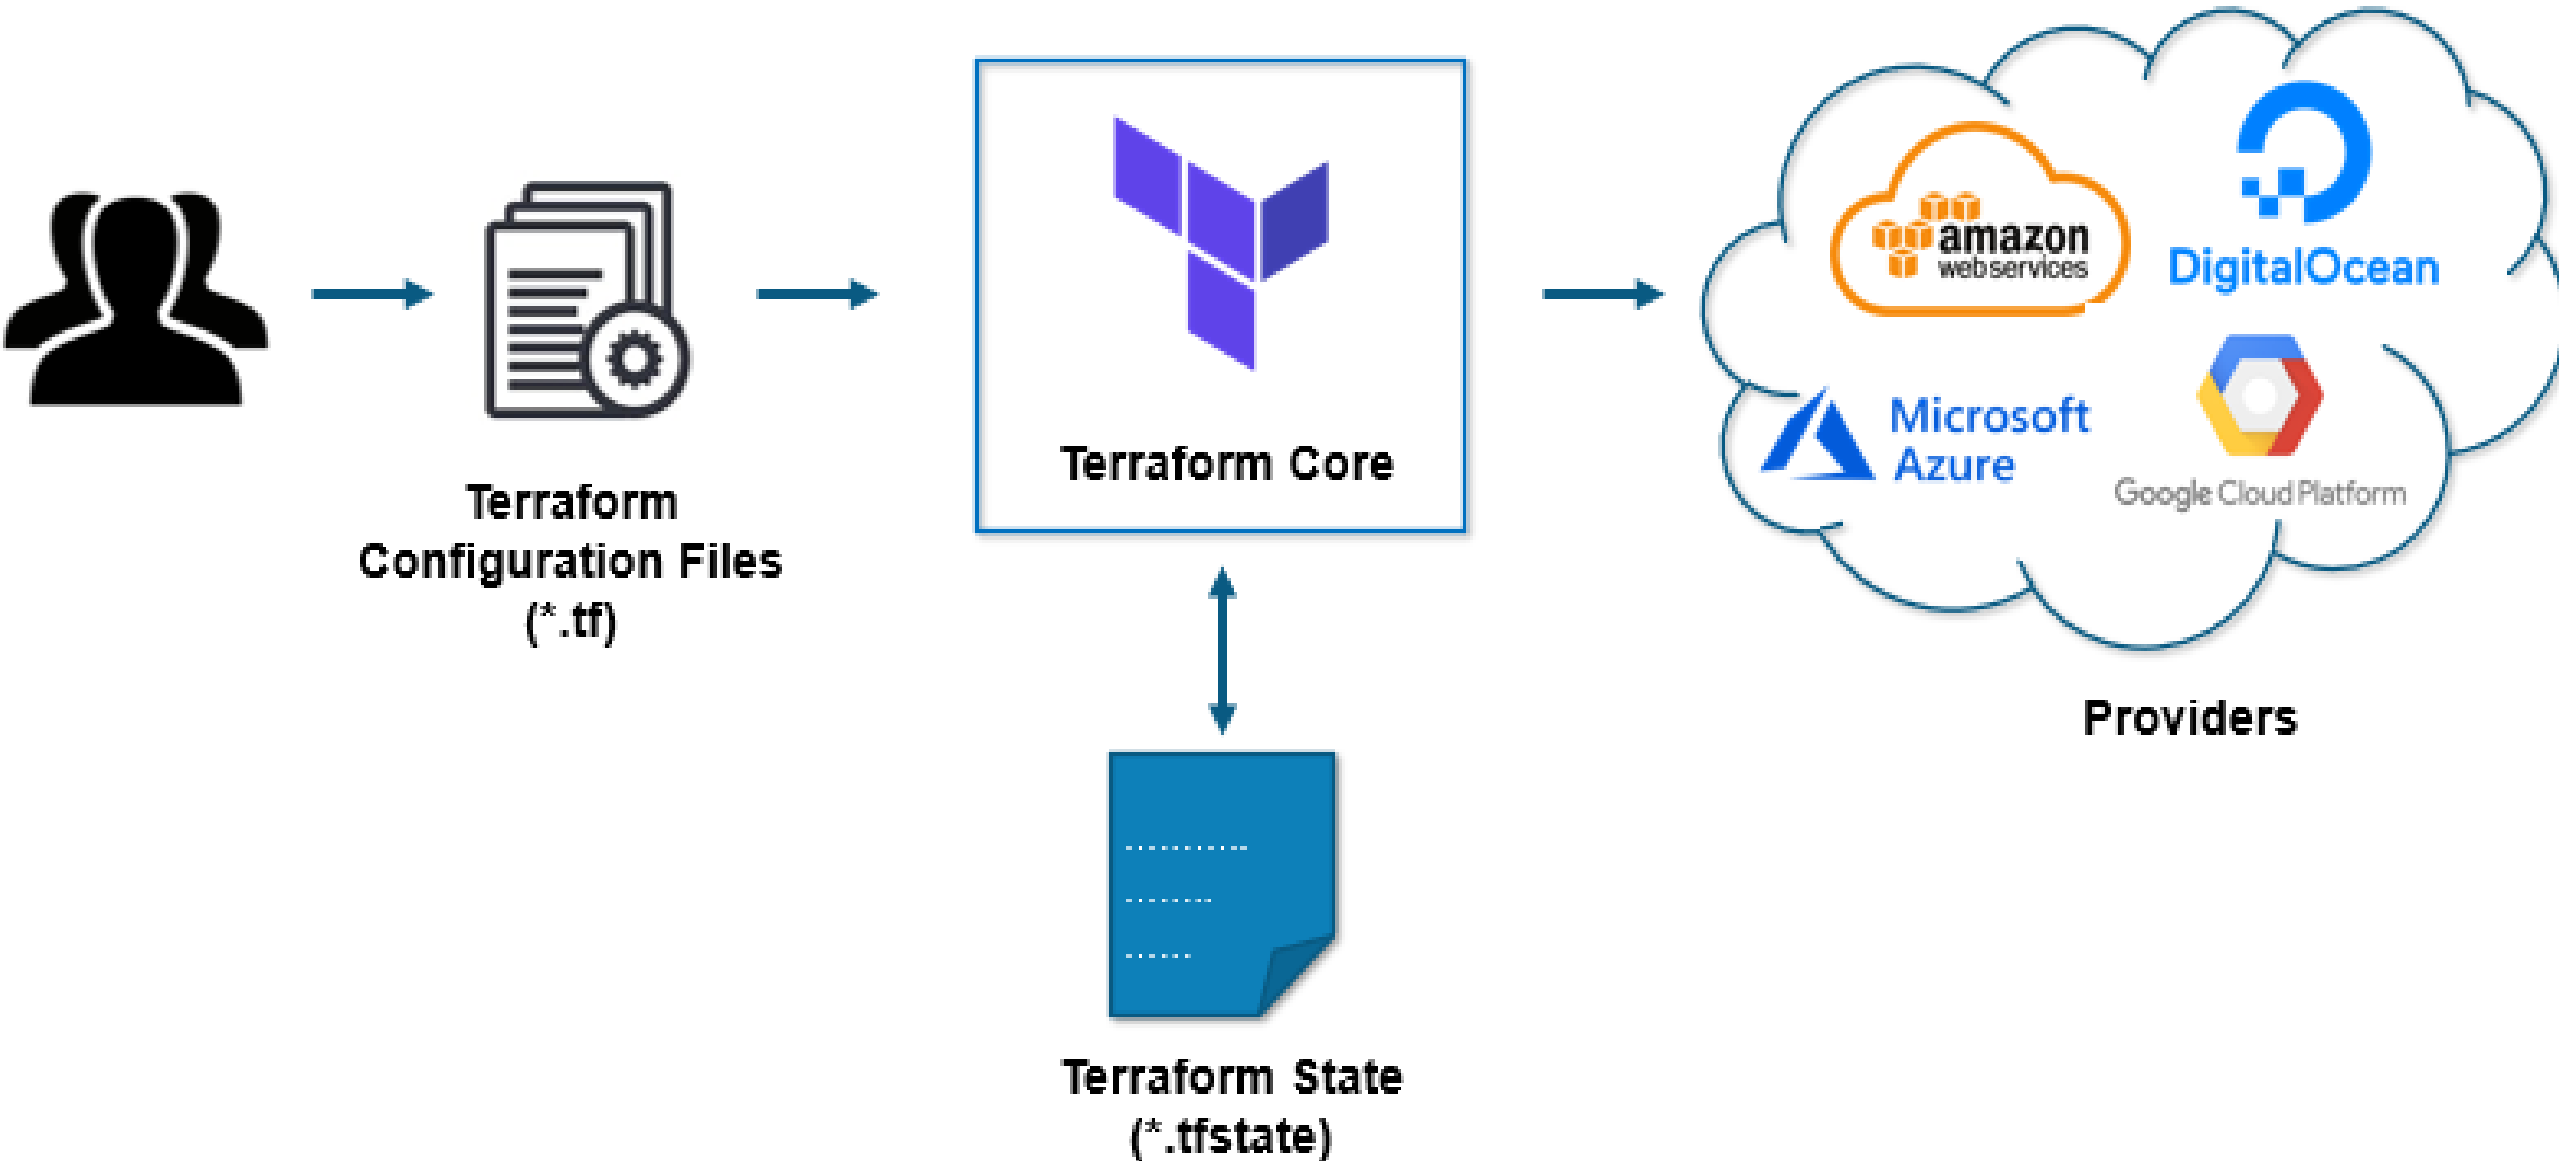
For example:

```
resource "aws_vpc" "main" {  
  cidr_block = var.base_cidr_block  
}
```

```
resource "local_file" "pet" {  
  filename = "filename_path"  
  content = "we love pets"  
}
```

Blocks are containers for other content and usually represent the configuration of some kind of object, like a resource.

Within the block body (between { and }) are the configuration arguments for the resource itself. Most arguments in this section depend on the resource type



Providers...



- **plugins called "providers" to interact with remote systems.**
 - Terraform relies on plugins called "providers" to interact with remote systems. Terraform configurations must declare which providers they require so that Terraform can install and use them. Additionally, some providers require configuration (like endpoint URLs or cloud regions) before they can be used.
- **What providers do?**
 - Each provider adds a set of resource types and/or data sources that Terraform can manage.
 - Every resource type is implemented by a provider; without providers, Terraform can't manage any kind of infrastructure.
 - Most providers configure a specific infrastructure platform (either cloud or self-hosted). Providers can also offer local utilities for tasks like generating random numbers for unique resource names.
- **Where Providers Come From**
 - Providers are distributed separately from Terraform itself, and each provider has its own release cadence and version numbers.
 - The Terraform Registry is the main directory of publicly available Terraform providers, and hosts providers for most major infrastructure platforms.


Providers...


Official


Verified

Community

[Browse Providers | Terraform Registry](#)

 HashiCorp Terraform | Registry


 Providers


 Modules

FILTERS

Clear Filters

Tier ?

☒  Official

☒  Verified

☒ Community

Category

☐ HashiCorp Platform

☐ Public Cloud

☐ Asset Management

☐ Cloud Automation

☐ Communication & Messaging

☐ Container Orchestration

☐ Continuous Integration/Deployment (CI/CD)

☐ Data Management

☐ Database

HOL



Create resource in aws

```
provider "aws" {  
  region = "ap-southeast-1"  
  access_key = ""  
  secret_key = ""  
}
```

```
resource "aws_instance" "ec2" { # I called name "ec2", you can change your own name  
  ami = "ami-0dad20bd1b9c8c004" # Image: Ubuntu Server 18.04 LTS  
  instance_type = "t2.micro" # VM Spec  
  security_groups = ["${aws_security_group.allow_ssh.name}"]  
  key_name = "aws-existingkey"  
}
```

```
resource "aws_security_group" "allow_ssh" {  
  name = "allow ssh"  
  ....[truncated]
```



Step Guide – AWS EC2 instance creation...

1. AWS cli installed system
2. Create AWS credential – programmatic access - portal
3. Add terraform block in main.tf file
4. Add provider block in main.tf file with authentication
5. Add resource block with required attributes

. init



> _

```
$ terraform init
```

```
Initializing the backend...
```

```
Initializing provider plugins...
```

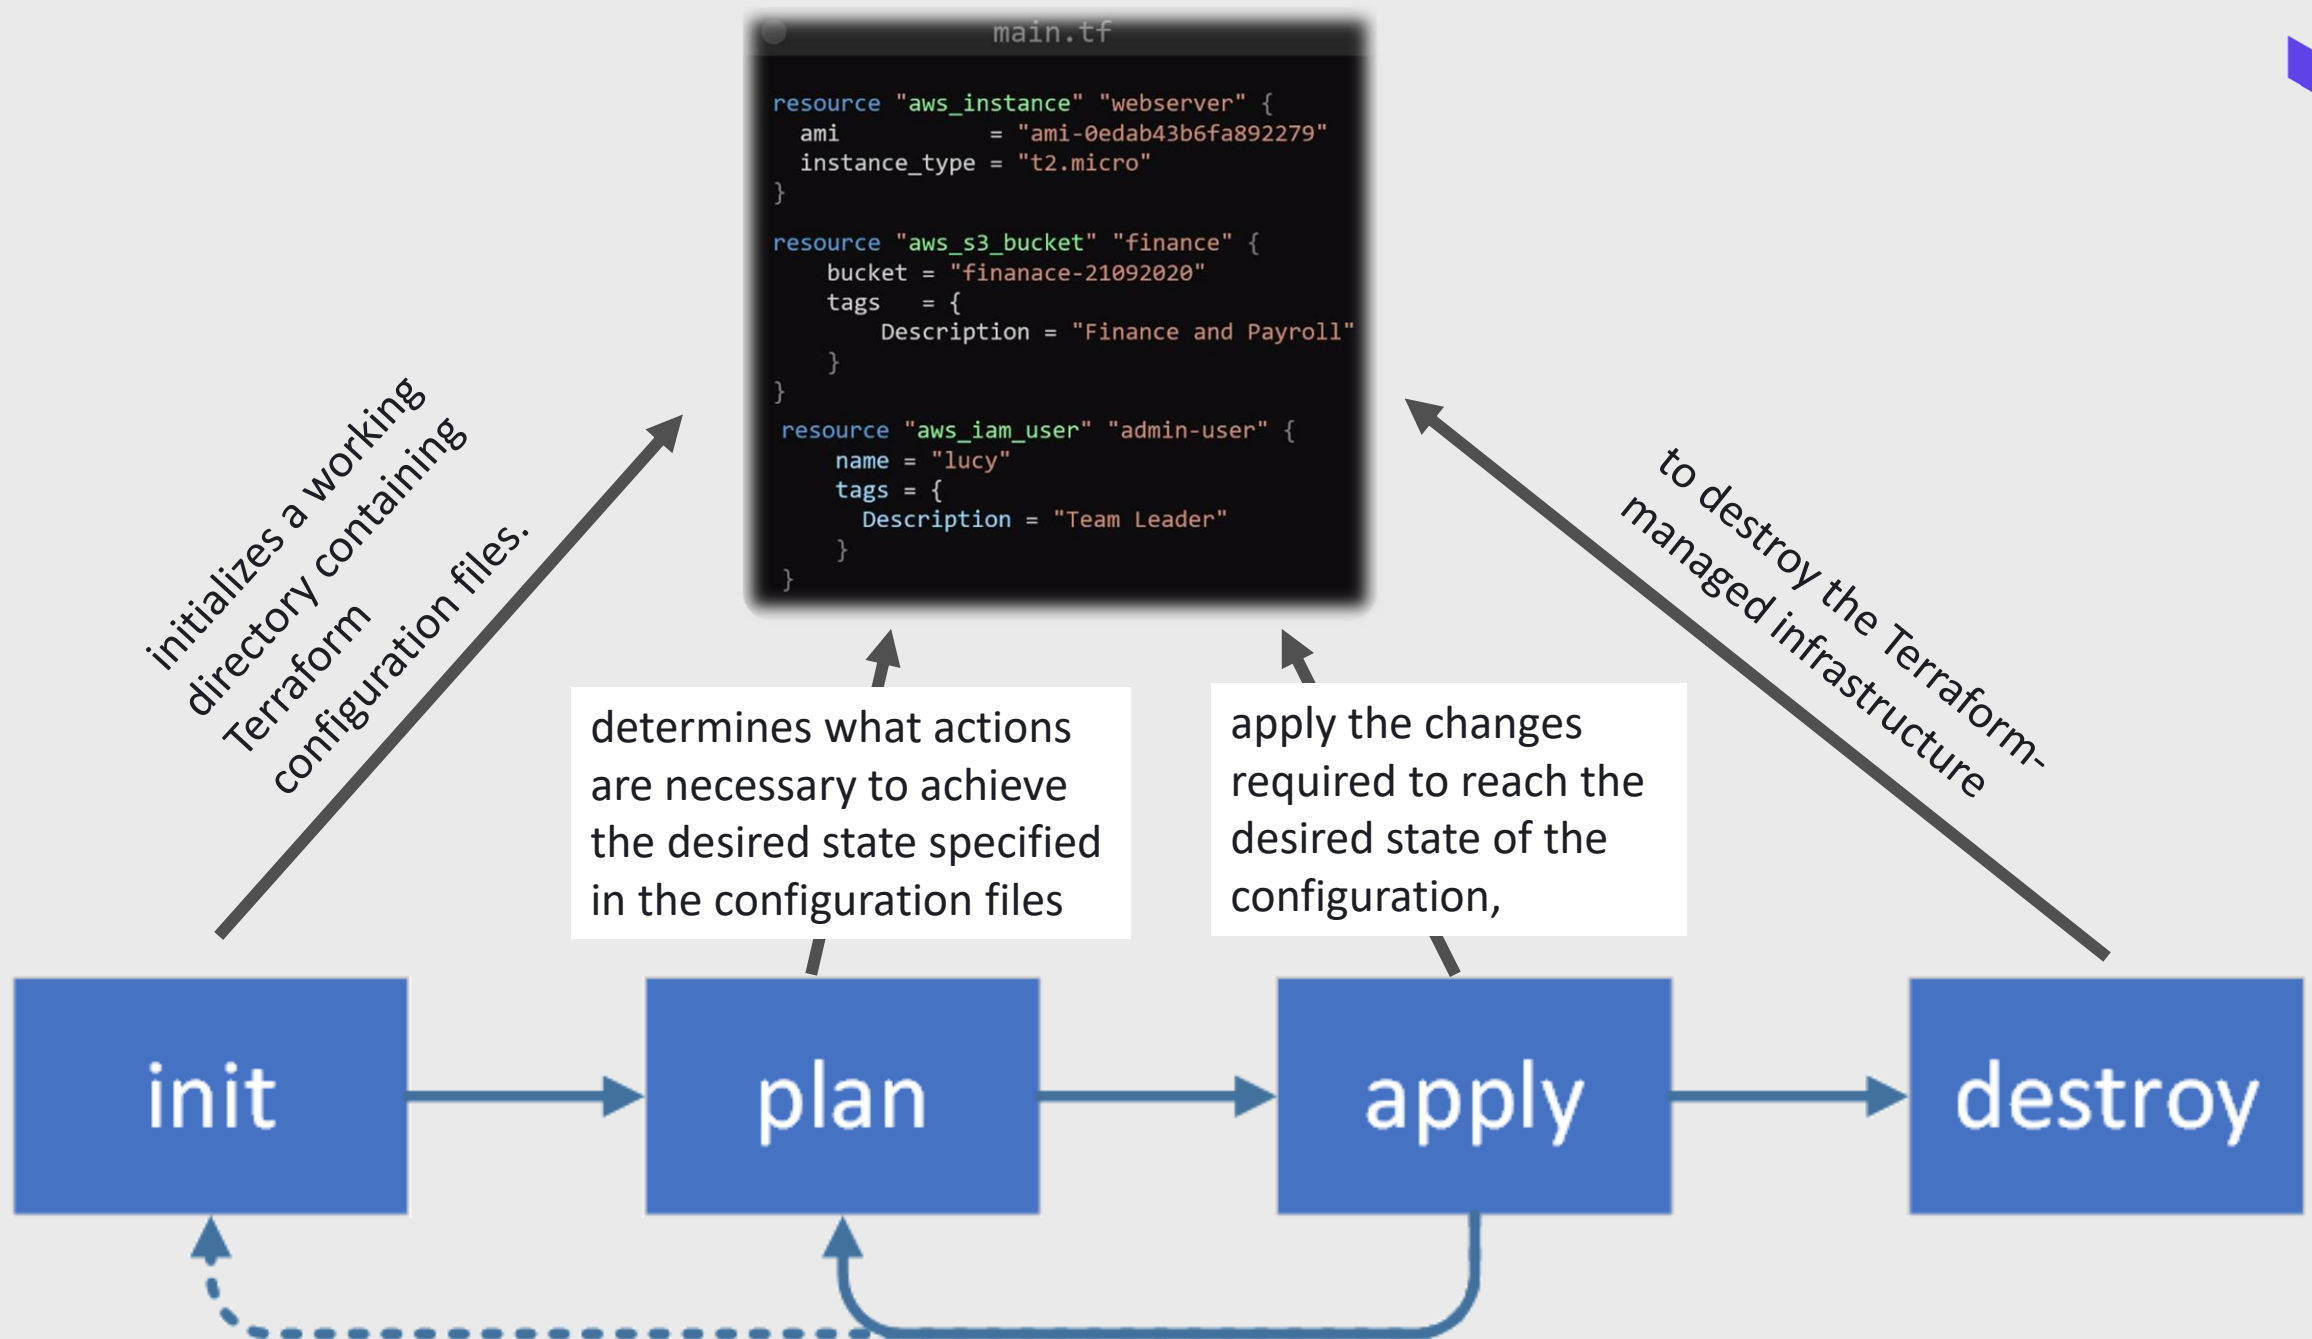
- Finding latest version of hashicorp/local...
- Installing hashicorp/local v2.0.0...
- Installed hashicorp/local v2.0.0 (signed by HashiCorp)

The following providers do not have any version constraints in configuration,
so the latest version was installed.

To prevent automatic upgrades to new major versions that may contain breaking changes, we recommend adding version constraints in a `required_providers` block in your configuration, with the constraint strings suggested below.

```
* hashicorp/local: version = "~> 2.0.0"
```

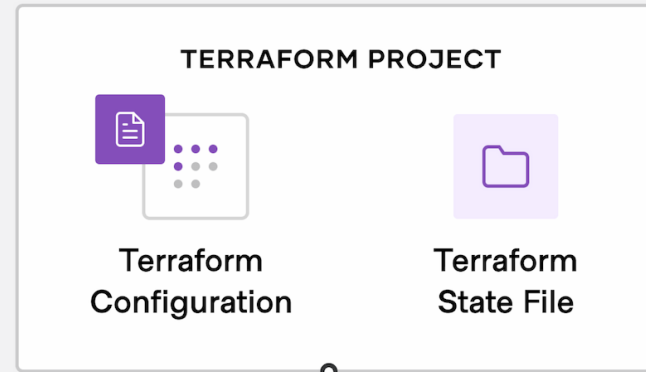
```
Terraform has been successfully initialized!
```





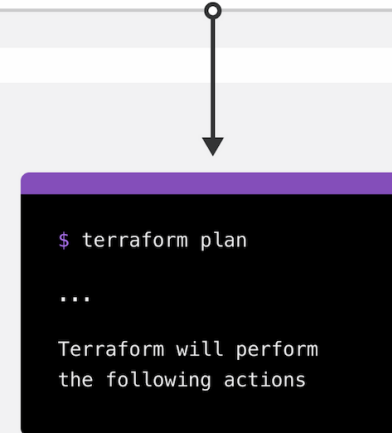
Write

Define infrastructure in configuration files



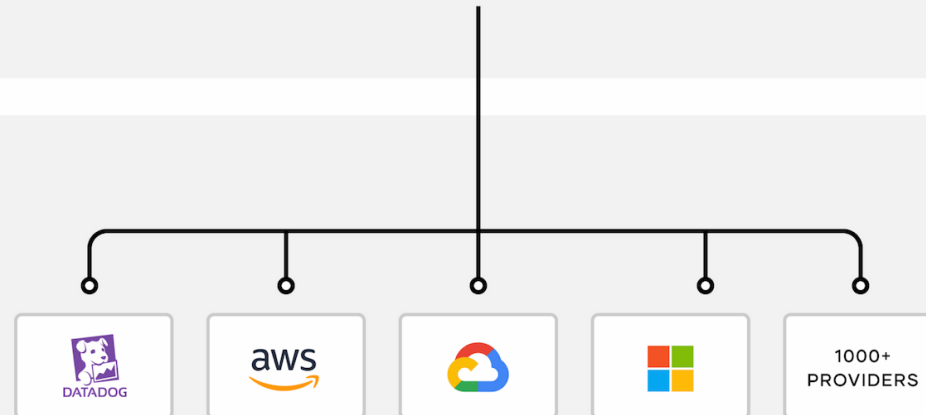
Plan

Review the changes
Terraform will make to
your infrastructure



Apply

Terraform provisions
your infrastructure and
updates the state file.





Variables..

- Input variables
 - used to define values that configure your infrastructure. These values can be used again and again without having to remember their every occurrence in the event it needs to be updated.
- Output variables
 - used to get information about the infrastructure after deployment. These can be useful for passing on information such as IP addresses for connecting to the server.



Input variables

Input variables are usually defined by stating a name, type and a default value.

The type and default values are not strictly necessary.

Terraform can deduct the type of the variable from the default or input value.

```
variable "variable_name" { }
```

```
variable "region" { }  
  
# Configure the AWS Provider  
provider "aws" {  
    region = "${var.region}"  
[...] truncated
```

```
terraform apply #will prompt region  
during execution
```

```
variable "region" {  
    default = "us-east-1"  
}  
  
# Configure the AWS Provider  
provider "aws" {  
    region = "${var.region}"  
[...] truncated
```

```
terraform apply #without supplying  
variable
```

```
terraform apply -var region="us-east-2"  
#with variable supplied...
```

Using variables in main.tf



main.tf

```
resource "local_file" "pet" {
  filename = var.filename
  content = var.content
}

resource "random_pet" "my-pet" {
  prefix = var.prefix
  separator = var.separator
  length = var.length
}
```

variables.tf

```
variable "filename" {
  default = "/root/pets.txt"
}

variable "content" {
  default = "We love pets!"
}

variable "prefix" {
  default = "Mrs"
}

variable "separator" {
  default = "."
}

variable "length" {
  default = "1"
}
```

In action..aws

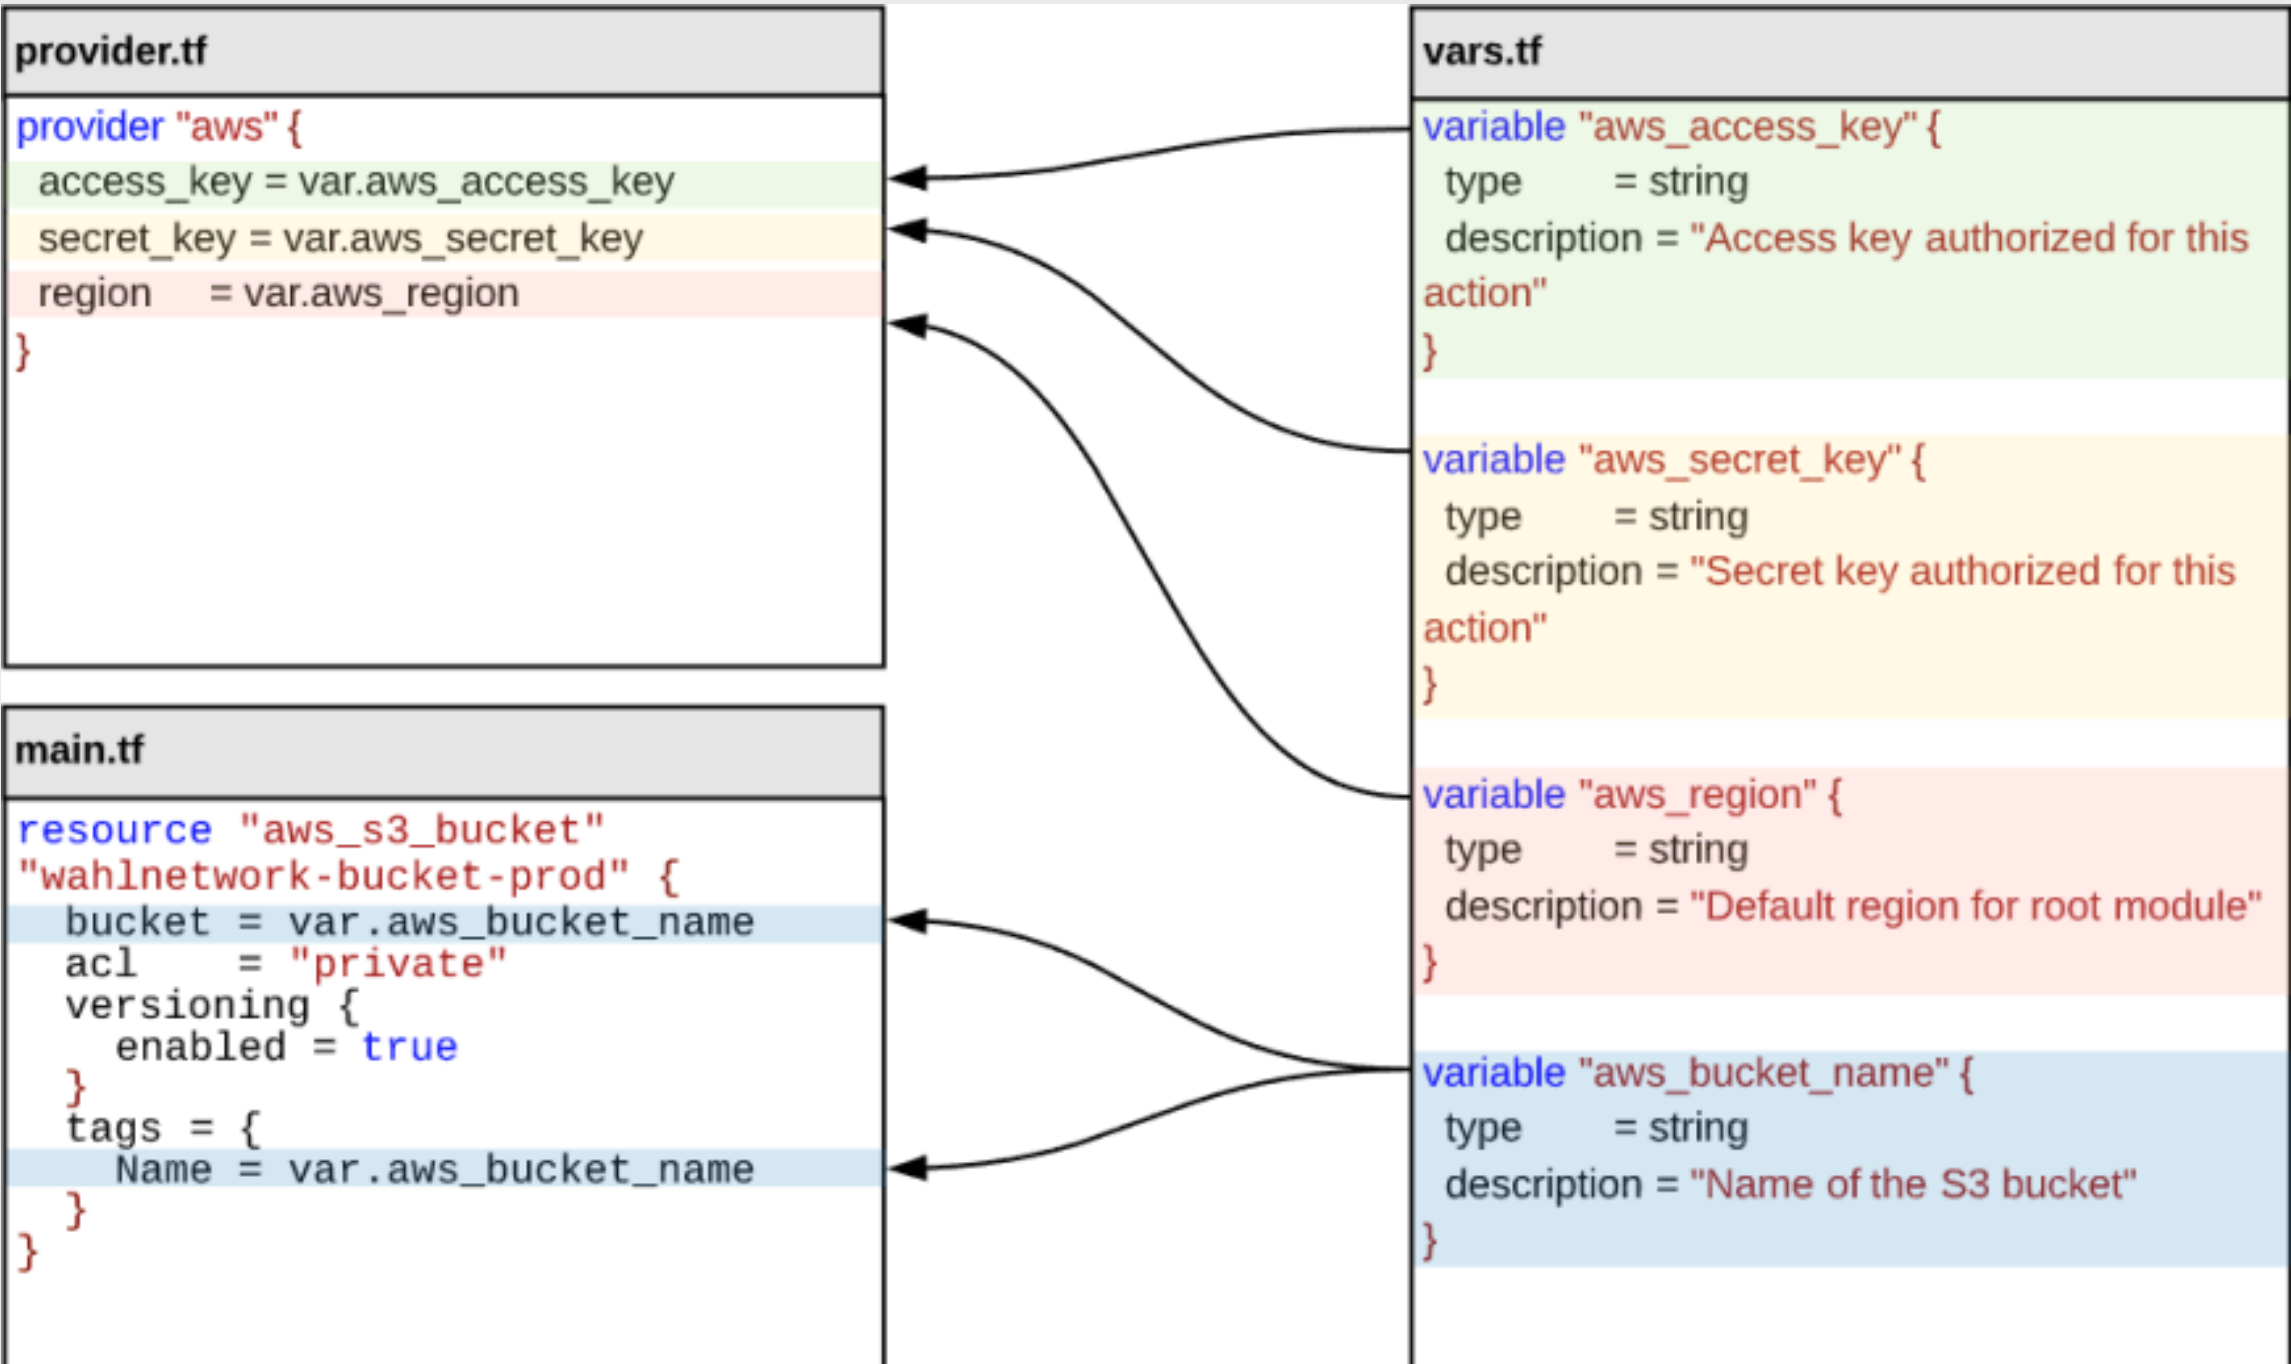


main.tf

```
resource "aws_instance" "webserver" {  
  ami           = var.ami  
  instance_type = var.instance_type  
}
```

variables.tf

```
variable "ami" {  
  default = "ami-0edab43b6fa892279"  
}  
variable "instance_type" {  
  default = "t2.micro"  
}
```





Output variables

provide a convenient way to get useful information about your infrastructure.

As you might have noticed, much of the resource details are calculated at deployment and only become available afterwards.

Using output variables you can extract any server-specific values including the calculated details.

```
terraform {...}
provider "aws" {...}
resource "aws_instance" "webserver" {...}
resource "local_file" "webserver_ip_details" {
  filename = "webserver_details.txt"
  content = <<-EOF
    public ip of webserver is ${aws_instance.webserver.public_ip}
    private ${aws_instance.webserver.private_ip}
  EOF
}
```

Config files...



File Name	Purpose
Main.tf	Contains resource definition. call modules, locals and data-sources to create all resources
Variables.tf	Contains variables declarations
Outputs.tf	Contains outputs from resources
Provider.tf	Contains provider definition



```
resource "aws_instance" "web" {  
  # ...
```


```
  provisioner "local-exec" {  
    command = "echo The server's IP address is ${self.private_ip}"  
  }  
}
```

```
resource "aws_instance" "webserver" {  
  ami = "ami-....."  
  instance_type = "t2.micro"  
  provisioner "remote-exec" {  
    inline = [  
      "yum -y install httpd",  
      "systemctl enable httpd",  
      "systemctl start httpd"  
    ]  
  }  
}
```

Modules....


[Standard Structure]


This is considered the root module and this is how you normally deploy infrastructure.



[Module Structure]

This is referred to as the child module as it is being called from the standard root module







```
└─ modules
    ├── ec2
    │   ├── main.tf
    │   ├── outputs.tf
    │   └── variables.tf
    └── vpc
        ├── main.tf
        ├── modules
        │   └── routing.tf
        │       └── subnets.tf
        ├── outputs.tf
        └── variables.tf
```




Tf modules

What are modules for?

- **Organize configuration** - Modules make it easier to navigate, understand, and update your configuration by keeping related parts of your configuration together. Even moderately complex infrastructure can require hundreds or thousands of lines of configuration to implement.
- **Encapsulate configuration** - Another benefit of using modules is to encapsulate configuration into distinct logical components. Encapsulation can help prevent unintended consequences, such as a change to one part of your configuration accidentally causing changes to other infrastructure, and reduce the chances of simple errors like using the same name for two different resources.
- **Re-use configuration** - Writing all of your configuration from scratch can be time consuming and error prone. Using modules can save time and reduce costly errors by re-using configuration written either by yourself, other members of your team, or other Terraform practitioners who have published modules for you to use. You can also share modules that you have written with your team or the general public, giving them the benefit of your hard work.
- **Provide consistency and ensure best practices** - Modules also help to provide consistency in your configurations. Not only does consistency make complex configurations easier to understand, it also helps to ensure that best practices are applied across all of your configuration. For instance, cloud providers give many options for configuring object storage services, such as Amazon S3 or Google Cloud Storage buckets. There have been many high-profile security incidents involving incorrectly secured object storage, and given the number of complex configuration options involved, it's easy to accidentally misconfigure these services.



Module

```
provider.tf

provider "aws" {
  access_key = var.aws_access_key
  secret_key = var.aws_secret_key
  region     = var.aws_region
}
```

```
main.tf

resource "aws_s3_bucket"
"wahlnetwork-bucket-prod" {
  bucket = var.aws_bucket_name
  acl     = "private"
  versioning {
    enabled = true
  }
  tags = {
    Name = var.aws_bucket_name
  }
}
```

```
vars.tf

variable "aws_access_key" {
  type      = string
  description = "Access key authorized for this action"
}

variable "aws_secret_key" {
  type      = string
  description = "Secret key authorized for this action"
}

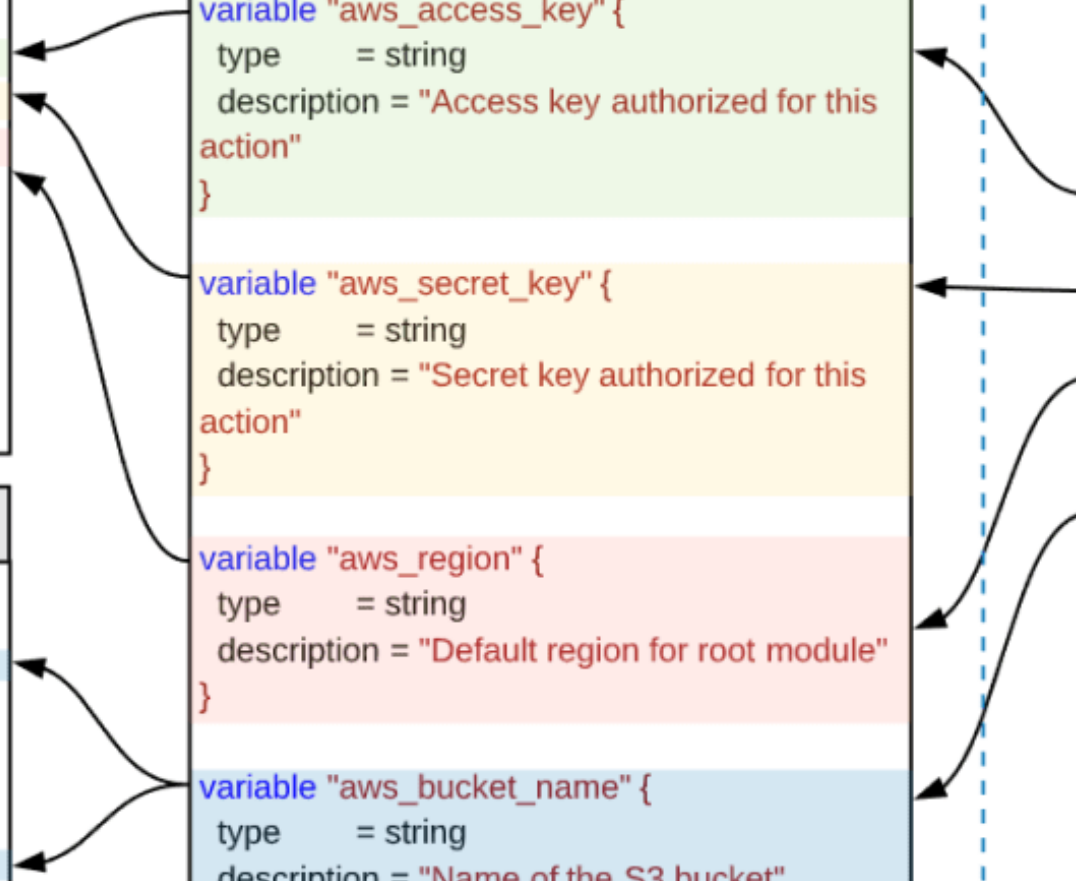
variable "aws_region" {
  type      = string
  description = "Default region for root module"
}

variable "aws_bucket_name" {
  type      = string
  description = "Name of the S3 bucket"
}
```

```
main.tf (dev)

module "s3-bucket" {
  source = "../../source"

  aws_access_key = "1234567890"
  aws_secret_key = "ABCDEFGHJIJ"
  aws_region     = "us-east-1"
  aws_bucket_name =
"wahlnetwork-bucket-dev"
}
```





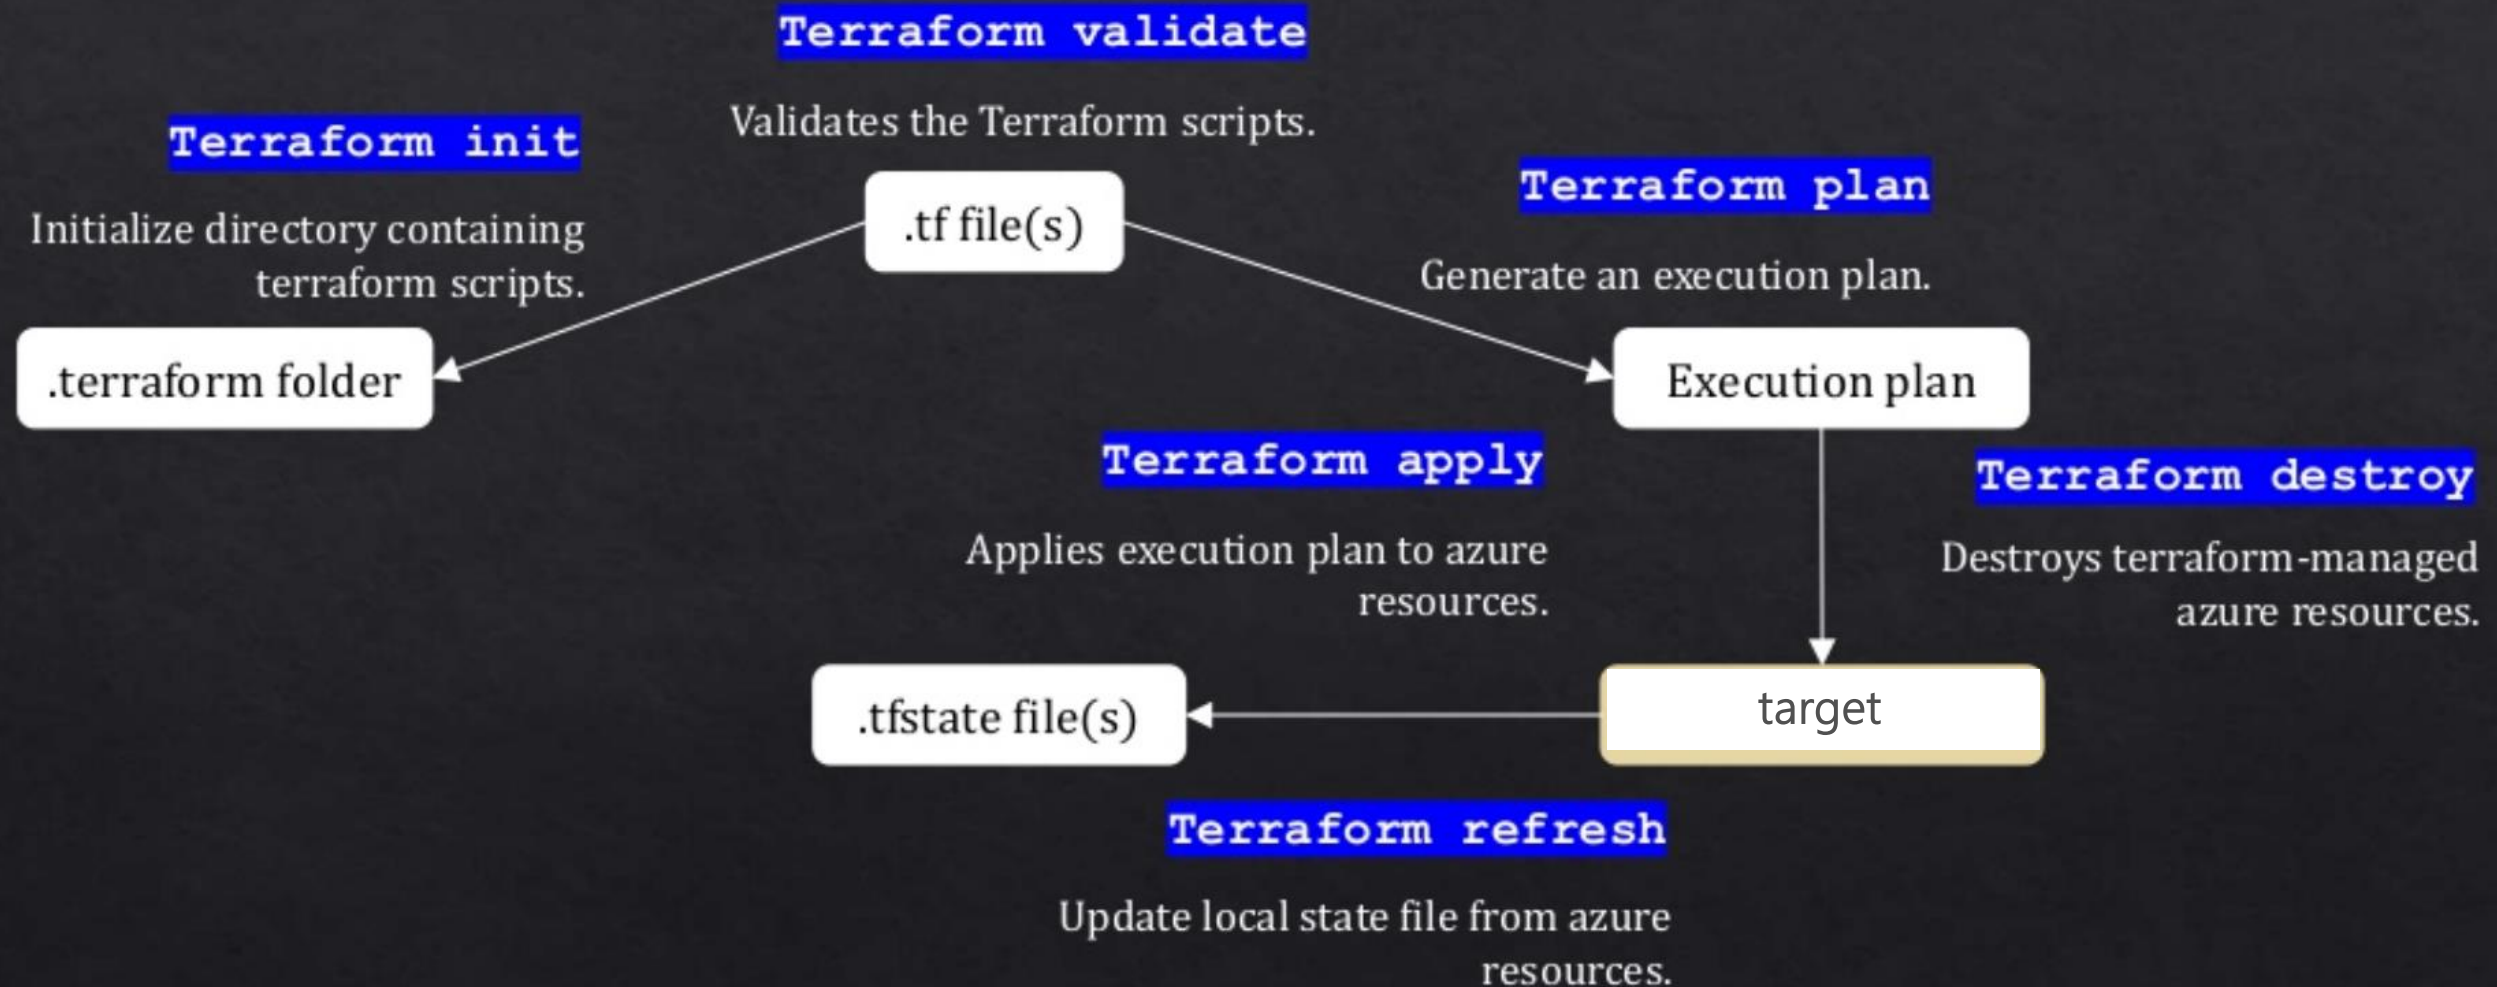
Local – with path

```
module "vpc_example_complete-vpc" {  
  source = "../modules/ec2instance"  
  version = "2.70.0"  
}
```

Tfregistry - <https://registry.terraform.io/browse/modules>

```
module "vpc_example_complete-vpc" {  
  source = "terraform-aws-modules/vpc/aws//examples/complete-vpc"  
  version = "2.70.0"  
}
```

tfState...



Tf Remote state



Version Control



Remote State Backends



main.tf

```
resource "local_file" "pet" {
  filename = "/root/pet.txt"
  content  = "My favorite pet is Mr.Whiskers!"
}
resource "random_pet" "my-pet" {
  length = 1
}
resource "local_file" "cat" {
  filename = "/root/cat.txt"
  content  = "I like cats too!"
}
```

terraform.tfstate

```
{
  "mode": "managed",
  "type": "aws_instance",
  "name": "dev-ec2",
  "provider": "provider[\"registry.terraform.io/hashicorp/aws\"]",
  "instances": [
    {
      "schema_version": 1,
      "attributes": {
        "ami": "ami-0a634ae95e11c6f91",
        .
        .
        .
        "primary_network_interface_id": "eni-0ccd57b1597e633e0",
        "private_dns": "ip-172-31-7-21.us-west-2.compute.internal",
        "private_ip": "172.31.7.21",
        "public_dns": "ec2-54-71-34-19.us-west-2.compute.amazonaws.com",
        "public_ip": "54.71.34.19",
        "root_block_device": [
          {
            "delete_on_termination": true,
            "device_name": "/dev/sda1",
            "encrypted": false,
            "iops": 100,
            "kms_key_id": "",
            "volume_id": "vol-070720a3636979c22",
            "volume_size": 8
          }
        ]
      }
    }
  ]
}
```

Terraform design flow

